

STATS 507

Data Analysis in Python

Week6-2: SciPy and Matplotlib


Dr. Xian Zhang

Adapted from slides by Professor Jeffrey Regier

Part 2: SciPy



1. SciPy as Scientific Python
2. Matplotlib for visualization



What is SciPy

SciPy: **Sci**entific **Py**thon

Designed and used for scientific and technical computing

SciPy is a library (collection) of **algorithms** and **mathematical** tools built to work with NumPy arrays.

- Mathematics, science, and engineering.
- Can handle operations and data analysis that goes beyond what basic Python or even NumPy can handle.

<https://docs.scipy.org/doc/scipy/tutorial/index.html#user-guide>

When to use SciPy

SciPy is organized into **subpackages** covering different scientific computing domains.

Special functions (`scipy.special`)

Integration (`scipy.integrate`)

Optimization (`scipy.optimize`)

Interpolation (`scipy.interpolate`)

Fourier Transforms (`scipy.fft`)

Signal Processing (`scipy.signal`)

Linear Algebra (`scipy.linalg`)

Sparse Arrays (`scipy.sparse`)

Sparse eigenvalue problems with ARPACK

Compressed Sparse Graph Routines (`scipy.sparse.csgraph`)

Spatial data structures and algorithms (`scipy.spatial`)

Statistics (`scipy.stats`)

Multidimensional image processing (`scipy.ndimage`)

File IO (`scipy.io`)

How to use SciPy: importing

Everything in the namespaces of SciPy submodules is **public**. In general in Python, it is recommended to make use of namespaces. For example:

```
my_matrix = np.array([
    [1, 2, 3],
    [4, 5, 6],
    [7, 8, 9]
])
# Method 1: Import the entire SciPy package
import scipy
# Usage
eigenvalues = scipy.linalg.eigvals(matrix)
```

```
# Method 2: Import specific submodules
from scipy import special, linalg
# Usage
eigenvalues = linalg.eigvals(matrix)
```

```
# Method 3: Import specific functions
from scipy.linalg import eigvals
# Usage
eigenvalues = eigvals(matrix)
```

```
print(type(eigenvalues))
print(eigenvalues)

<class 'numpy.ndarray'>
[ 1.61168440e+01+0.j -1.11684397e+00+0.j -8.58274334e-16+0.j]
```

SciPy Linear Algebra

There is overlap in the functionality provided by the SciPy and NumPy submodules. However:

- Slightly different from `numpy.linalg`. Always uses BLAS/LAPACK support, which are often highly optimized for specific hardware so could be **faster**.
- Unless you don't want to add SciPy as a dependency to your NumPy program, use `scipy.linalg` instead of `numpy.linalg`

BLAS: basic linear algebra subprogram: <https://www.netlib.org/blas/>

LAPACK: linear algebra Package: <https://www.netlib.org/lapack>

<https://docs.scipy.org/doc/scipy/tutorial/linalg.html>

SciPy Linear Algebra

`scipy.linalg` contains all the functions in `numpy.linalg`

SciPy linear algebra could be **faster**.

<https://docs.scipy.org/doc/scipy/tutorial/linalg.html>

```
import numpy as np
import scipy.linalg as spla
import numpy.linalg as npla
import time
```

Solve a linear system: $Ax = b$

```
# Create a large matrix
```

```
n = 1000
```

```
A = np.random.rand(n, n)
```

```
b = np.random.rand(n)
```

```
# NumPy solve
```

```
start_time = time.time()
```

```
x_np = npla.solve(A, b)
```

```
np_time = time.time() - start_time
```

```
# SciPy solve
```

```
start_time = time.time()
```

```
x_sp = spla.solve(A, b)
```

```
sp_time = time.time() - start_time
```

```
print(f"NumPy solve time: {np_time:.6f} seconds")
```

```
print(f"SciPy solve time: {sp_time:.6f} seconds")
```

```
print(f"Speed difference: {np_time/sp_time:.2f}x")
```

```
NumPy solve time: 0.044730 seconds
```

```
SciPy solve time: 0.020878 seconds
```

```
Speed difference: 2.14x
```


SciPy Sparse

Sparse matrix classes: CSC, CSR, etc.

- functions to **identify** and **build** sparse matrices
- `sparse.linalg` module for sparse linear algebra
- `sparse.csgraph` for sparse graph routines

```
import numpy as np
from scipy import sparse
# 1. Sparse matrix classes: CSC, CSR
dense_matrix = np.array([
    [1, 0, 2],
    [0, 3, 4],
    [5, 6, 0]
])

# Create Compressed Sparse row matrix
csr_matrix = sparse.csr_matrix(dense_matrix)
print("CSR Matrix:")
print(type(csr_matrix))
print(csr_matrix)

CSR Matrix:
<class 'scipy.sparse._csr.csr_matrix'>
(0, 0)    1
(0, 2)    2
(1, 1)    3
(1, 2)    4
(2, 0)    5
(2, 1)    6
```

```
# sparse.linalg module for sparse linear algebra
# Solve a sparse linear system Ax = b
from scipy.sparse import linalg as spla
A = sparse.csr_matrix([[1, 2], [3, 0]])
b = np.array([1, 2])
x = spla.spsolve(A, b)
print("\nSolution to Ax = b:")
print(x)
```

Can be even more efficient
Since we can make use of
the structure

More subroutines

Scipy **Statistics**

- Distributions
- Functions (Mean, median, mode, variance)
- Hypothesis tests
- ...

More on this later!

Scipy **Signal**

- Continuous-time linear system
- Filtering
- ...

Scipy **IO**

- Methods for loading and saving data
- Matlab files
- Matrix Market files (sparse matrices)
- Wav files

In-class practice

On SciPy Optimization

Part 3: Matplotlib



What is matplotlib

Matplotlib is a comprehensive **library** for creating static, animated, and interactive **visualizations** in Python.

Similar to R's `ggplot2` and MATLAB's plotting functions

For MATLAB fans, matplotlib.pyplot implements MATLAB-like plotting:

http://matplotlib.org/users/pyplot_tutorial.html

Sample plots with code:


http://matplotlib.org/tutorials/introductory/sample_plots.html

Basic plotting: `matplotlib.pyplot.plot`

`matplotlib.pyplot.plot(x, y)`
plots y as a function of x .

`matplotlib.pyplot.plot(y)`
default x-axis to `np.arange(len(x))`

```
import numpy as np
import matplotlib.pyplot as plt
# Create some example data
x = np.linspace(0, 10, 100)
y = np.sin(x)
# Create the plot
plt.plot(x, y)
# Display the plot
plt.show()
```




`matplotlib.pyplot` is the main plotting interface in Matplotlib. It provides an implicit, MATLAB-like way of plotting, often imported as `plt`

Basic plotting: `matplotlib.pyplot.plot`

`matplotlib.pyplot.plot(x, y)`
plots y as a function of x .

```
import numpy as np
import matplotlib.pyplot as plt
# Create some example data
x = np.linspace(0, 10, 100)
y = np.sin(x)
# Create the plot
plt.plot(x, y)
# Add a title
plt.title("Simple Sine Wave Example")
# Add an axis label
plt.xlabel("X-axis")
plt.ylabel("Y-axis")
# Display the plot
plt.show()
```


Specifying titles and axis labels
couldn't be more straight-forward.



Customizing plots

```
1 x = np.arange(0,5,0.25, dtype='float')
2 _ = plt.plot(x**2, ':ro')
```

Second argument to `pyplot.plot` specifies line type, line color, and marker type.



matplotlib.pyplot.plot


matplotlib.pyplot.plot(*args, **kwargs)

positional argument

Optional keyword arguments
for **customizing** the plot

Customizing plots

```
1 x = np.arange(0,5,0.25, dtype='float')  
2 _ = plt.plot(x**2, color='red', linestyle=':', marker='o')
```




Long form of the command on the previous slide. Same plot!

format **string** characters are accepted to control the color, line style or marker:

A full list of the long-form arguments available to `pyplot.plot` are available in the table titled “Here are the available Line2D properties.”:
http://matplotlib.org/users/pyplot_tutorial.html

Multiple lines in a single plot

```
1 t = np.arange(0., 5., 0.2)
2 # plt.plot(xvals, yvals, traits1, y2vals, traits2, ... )
3 _ = plt.plot(t, t, 'r--', t, t**2, 'bs', t, t**3, 'g^')
```




Note: more complicated specification of individual lines can be achieved by adding them to the plot one at a time.

Multiple lines in a single plot

```
1 t = np.arange(0., 5., 0.2)
2 plt.grid()
3 plt.plot(t, t, 'r--')
4 plt.plot(t, t**2, 'bs')
5 plt.plot(t, t**3, 'g^')
6 _ = plt.show()
```

`plt.grid` to apply grid to the figure

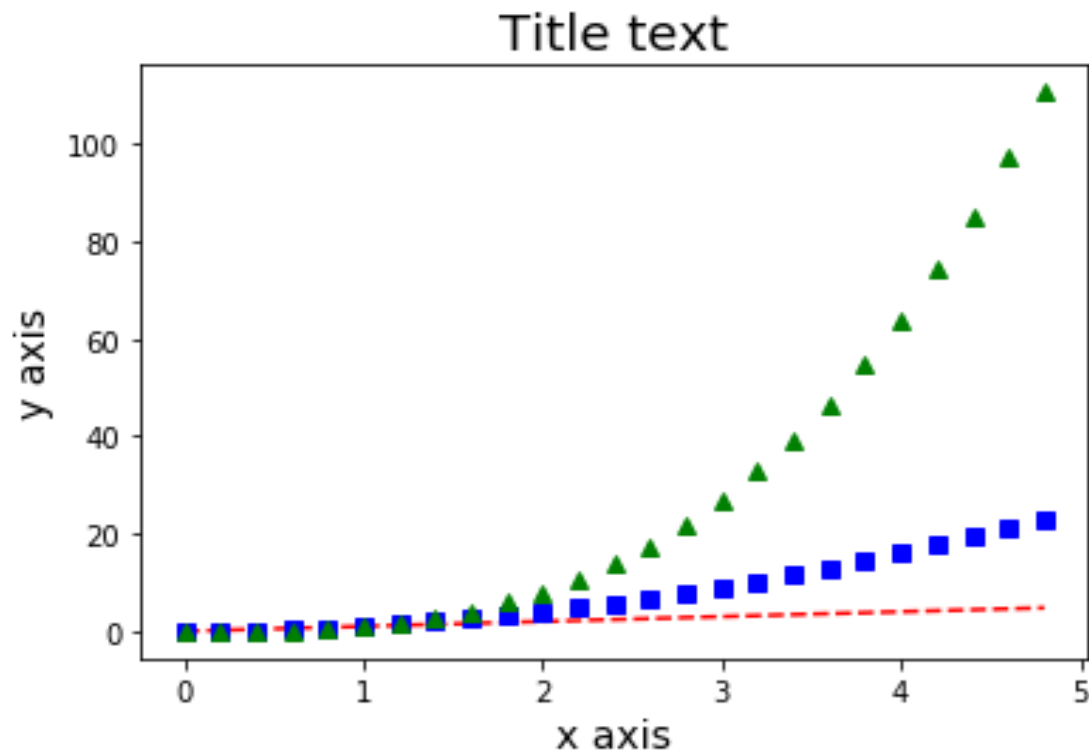


Note: same plot as previous slide, but specifying one line at a time.

Titles and axis labels

```
1 t = np.arange(0., 5., 0.2)
2 plt.title('Title text', fontsize=18)
3 plt.xlabel('x axis', fontsize=14)
4 plt.ylabel('y axis', fontsize=14)
5 _ = plt.plot(t, t, 'r--', t, t**2, 'bs', t, t**3, 'g^')
```


Change font sizes



Legends

```
1 plt.xlabel("$n$", fontsize=16) # set the axes labels
2 plt.ylabel("$f(n)$", fontsize=16)
3 plt.title("Different growth behaviors") # set the plot title
4 plt.plot(t, t, '-ob', label='linear, $f(n)=n$')
5 plt.plot(t, t**2, ':-^r', label='quadratic, $f(n)=n^2$')
6 plt.plot(t, t**3, '--sg', label='cubic, $f(n)=n^3$')
7 _ = plt.legend(loc='best') # places legend at best location
```


Can use LaTeX in labels, titles, etc.



`pyplot.legend` generates legend based on label arguments passed to `pyplot.plot`. `loc='best'` tells `pyplot` to place the legend where it thinks is best.

Annotating figures

```
1 t = np.arange(0.0, 5.0, 0.01)
2 s = np.cos(2*np.pi*t) #np.pi==3.14159...
3 plt.plot(t, s, lw=2) # plot the cosine.
4 # Annotate the figure with an arrow and text.
5 _ = plt.annotate('local max', xy=(2, 1), xytext=(3, 1.5),
6                 fontsize=14,
7                 arrowprops=dict(facecolor='black', shrink=0.02) )
```




Specify text coordinates and coordinates of the arrowhead using the *coordinates of the plot itself*. This is pleasantly different from many other plotting packages, which require specifying coordinates in pixels or inches/cms.

Plotting histograms: `pyplot.hist()`

Bin counts: defines the number of equal-width bins in the histogram or the bin edges.

```
1 mu, sigma = (100, 15)
2 x = np.random.normal(mu, sigma, 10000)
3 # hist( data, nbins, ... )
4 (n, bins, patches) = plt.hist(x, 50, density=False, facecolor='teal')
5 n
```

```
array([ 1.,  1.,  2.,  4.,  3.,  5., 11., 18., 26., 30., 47.,
       68., 82., 113., 150., 201., 246., 285., 309., 352., 420., 475.,
       541., 529., 597., 595., 572., 566., 543., 515., 462., 404., 360.,
       270., 294., 233., 159., 128., 111., 92., 54., 32., 28., 28.,
       15., 11., 5., 2., 1., 4.] )
```




Note that if `density=True`, then these will be chosen so that the histogram “integrates” to 1. It is normalized so the total area equals 1.

https://matplotlib.org/3.1.1/api/as_gen/matplotlib.pyplot.hist.html

Bar plots

```
1 import string
2 t = np.arange(10)
3 s = np.random.normal(1,1,10)
4 mylabels = list(string.ascii_lowercase[0:len(t)])
5 _ = plt.bar(t, s, tick_label=mylabels, align='center')
```



Full set of available arguments to `bar(...)` can be found at

http://matplotlib.org/api/_as_gen/matplotlib.pyplot.bar.html#matplotlib.pyplot.bar

Horizontal analogue given by `barh`


http://matplotlib.org/api/_as_gen/matplotlib.pyplot.barh.html#matplotlib.pyplot.barh

Can specify what the x-axis tick labels should be by using the `tick_label` argument to plot functions.

Box & whisker plots

Draw a box and whisker plot. The box extends from the first quartile (Q1: a statistical measure that represents the 25th percentile of a dataset.) to the third quartile (Q3) of the data, with a line at the median.

```
1 K=12; n=25
2 draws = np.zeros((n,K))
3 for k in range(K):
4     mu = np.sin(2*np.pi*k/K)
5     draws[:,k] = np.random.normal(mu,1,n)
6 _ = plt.boxplot(draws, labels=list('JFMAMJJASOND'))
```



`plt.boxplot(x, ...)` : `x` is the data. Many more optional arguments are available, most to do with how to compute medians, confidence intervals, whiskers, etc. See http://matplotlib.org/api/_as_gen/matplotlib.pyplot.boxplot.html#matplotlib.pyplot.boxplot

Pie Charts

Don't use pie charts!

A table is nearly always better than a dumb pie chart; the only worse design than a pie chart is several of them, for then the viewer is asked to compare quantities located in spatial disarray both within and between charts [...]
Given their low [information] density and failure to order numbers along a visual dimension, pie charts should never be used.

Edward Tufte

The Visual Display of Quantitative Information

But if you must...

```
pyplot.pie(x, ... )
```

http://matplotlib.org/api/_as_gen/matplotlib.pyplot.pie.html#matplotlib.pyplot.pie



Subplots

`subplot(nrows, ncols, plot_number)`

Shorthand: `subplot(XYZ)`

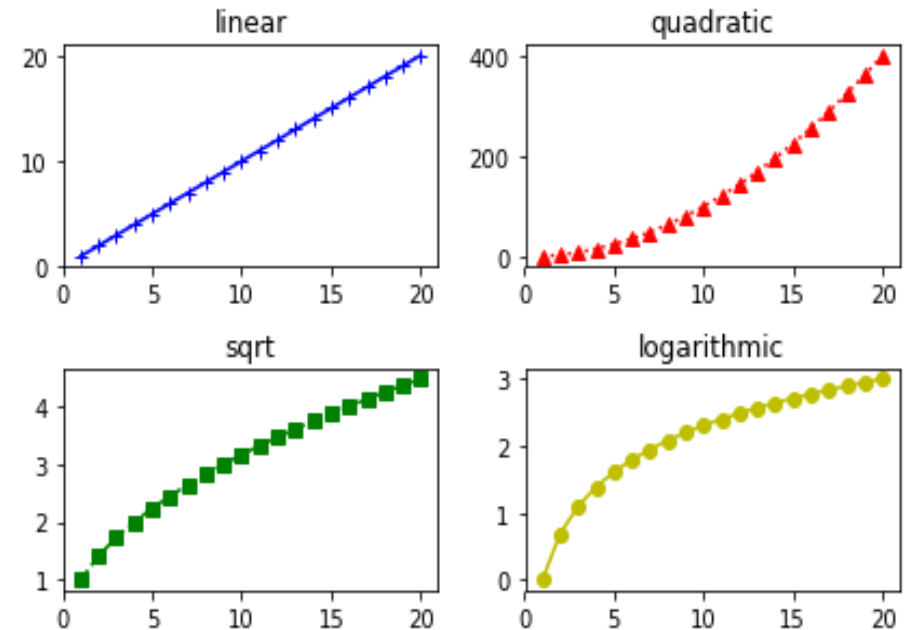
Makes an X-by-Y plot

Picks out the Z-th plot

Counting in row-major order

`tight_layout()` automatically tries to clean things up so that subplots don't overlap. Without this command in this example, the labels "sqrt" and "logarithmic" overlap with the x-axis tick labels in the first row.

```
1 t=np.arange(20)+1
2 plt.subplot(221)
3 plt.plot(t,t,'-+b')
4 plt.title('linear')
5 plt.subplot(222)
6 plt.title('quadratic')
7 plt.plot(t,t**2,':^r')
8 plt.subplot(223)
9 plt.title('sqrt')
10 plt.plot(t,np.sqrt(t),'--sg')
11 plt.subplot(224)
12 plt.title('logarithmic')
13 plt.plot(t,np.log(t),'-oy')
14 _ = plt.tight_layout()
```



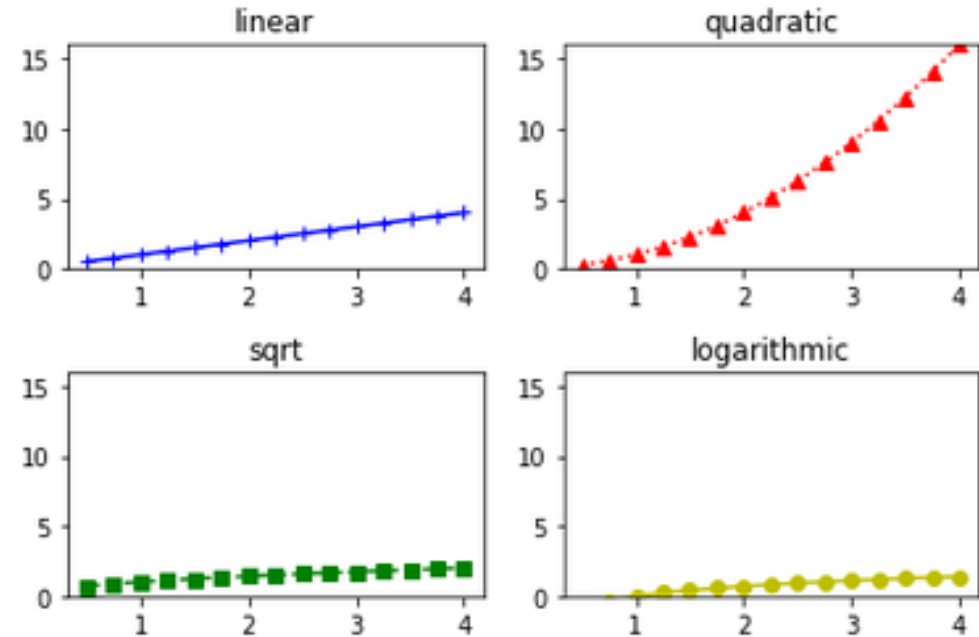
Specifying axis ranges

`plt.ylim([lower, upper])` sets y-axis limits

`plt.xlim([lower, upper])` for x-axis

For-loop goes through all of the subplots and sets their y-axis limits

```
1 t = np.arange(0.5, 4.25, 0.25)
2 ymax = np.max(t**2)
3 plt.subplot(221)
4 plt.plot(t, t, '-+b')
5 plt.title('linear')
6 plt.subplot(222)
7 plt.title('quadratic')
8 plt.plot(t, t**2, 'r')
9 plt.subplot(223)
10 plt.title('sqrt')
11 plt.plot(t, np.sqrt(t), '--sg')
12 plt.subplot(224)
13 plt.title('logarithmic')
14 plt.plot(t, np.log(t), '-oy')
15 for subplt in range(221, 225):
16     plt.subplot(subplt)
17     plt.ylim([0, ymax])
18 _ = plt.tight_layout()
```



Nonlinear axis

Scale the axes with

`plt.xscale` and `plt.yscale`

Built-in scales:

- Linear ('linear')

- Log ('log')


- Symmetric log ('symlog')

- Logit ('logit')

Can also specify customized scales:

https://matplotlib.org/devel/add_new_projection.html#adding-new-scales

```
1 y = np.random.uniform(0,1,10000); y.sort()
2 x = np.arange(len(y))
3 plt.subplot(121)
4 plt.plot(x,y)
5 plt.yscale('linear'); plt.title('linear')
6 plt.subplot(122)
7 plt.plot(x, y)
8 plt.yscale('log'); plt.title('log')
9 _ = plt.tight_layout()
```



Saving images

`plt.savefig(filename)` will try to **automatically** figure out what file type you want based on the file extension.

Or Can make it explicit using

`plt.savefig('filename',
format='fmt')`

Options for specifying resolution, padding, etc:

https://matplotlib.org/api/_as_gen/matplotlib.pyplot.savefig.html

```
1 random_signs = np.sign(np.random.rand(1000)-0.5)
2 plt.grid(True)
3 plt.title('Random walk of 1000 steps')
4 # cumsum() returns cumulative sums
5 _ = plt.plot(np.cumsum(random_signs))
6 plt.savefig('random_walk.svg')
```



Animations

`matplotlib.animate` package generates animations

We won't require you to make any, but they're fun to play around with (and they can be a great visualization tool)

The details are a bit tricky, so I recommend starting by looking at some of the example animations here:

https://matplotlib.org/stable/api/animation_api.html

In-class practice

Other things

HW5 due Today.

HW6 is out and due after Fall Break.

Coming next:

- Midterm recap

- Midterm